

Application-Level Energy Accounting with Chappie – Technical Report

Anonymous Author(s)

1 Overview

This document provides supplemental material for the CHAPPIE project. Included are a technical report on our experiences working with CHAPPIE inside a simulator, and experimental results.

2 Energy Simulation

In this section, we document our (unsuccessful) effort in simulating per-method energy consumption through a cycle-accurate simulator. We faced with two challenges: (1) full-system simulation for realistic Java benchmarks is extremely time-consuming; (2) The behavior of an instrumented benchmark under simulation — where method entry/exit are instrumented with special instructions for demarcating the boundary of energy simulation — significantly differs from that of its unmodified counterpart.

2.1 Simulation Setup

We used gem5 version 2.0 [1] to simulate running a full Ubuntu 16.04 system. In part, we choose this approach to run the Java Virtual Machine (JVM) inside a cycle-accurate simulator. gem5 simulation statistics were piped through mcPAT version 1.3 [2] to simulate power numbers. Our experiments from CHAPPIE base system were performed on a dual socket Intel E5-2630 v4 2.20 GHz CPU server. Please see the main paper Section 5.3 for full details.

We structured two experiments designed to evaluate CHAPPIE in the context of simulation: (E1) we instrumented the top ranked methods for a selected benchmark—as determined by CHAPPIE method energy accounting running on the base system—with code to capture the simulation cycles of the method execution. Our goal was to compare these numbers with base system energy accounting numbers; (E2) we ran CHAPPIE for the selected benchmark on the simulator—we modified CHAPPIE to read simulator cycles instead of energy registers and convert the simulator cycles to energy a posteriori with mcPAT—and compare the energy recorded in (E1) with that attributed by CHAPPIE. Our hope was to cross validate the method energy number produced by the cycle-accurate simulate in (E1) with CHAPPIE energy accounting using the same cycle-accurate simulation.

2.2 Discussion

Our efforts revealed two fundamental problems with this approach: overhead, and more importantly, application behavior change.

| benchmark | base runtime | gem5 runtime |
|-----------|--------------|--------------|
| sunflow | 1.5 (s) | 3.29 (hr) |
| h2 | 11.4 (s) | DNF |

Figure 1. Benchmark Runtime Comparison for gem5

We show the runtime of our gem5 exploration experiments for sunflow and h2 in Figure 1. sunflow, the smaller of the two benchmarks, completed in 3.29 hours on the gem5 simulator, with a 7895x slowdown compared with the base system. h2 did not complete. To understand the difficulty with h2, we turned to our experiment to instrument one of the top methods, compareRow for h2, and found that the method is invoked for 58 million times. Based on our observations, we estimate a single run of this invocation on the simulator takes 1.5 seconds, and total execution time is likely to be 990 days. We are aware that speeding up cycle accurate simulation is an active research area [3].

Overhead, however, is the lesser of the problems for us. We implemented an early exit scheme that would collect data for an experiment until a particular top method (compareRows in this case) was invoked 100 times. For our experiment where we monitor with CHAPPIE and measure method invocation (E2) in the same run, the results showed that not only was the instrumented code slowing down the application—a manifestation of the scenario described in Section 2 of our main paper—but that the methods attributed energy to were different from the base system; namely, the instrumentation code was altering the application behavior.

Engineering Effort The gem5 simulation has two methods for simulating programs: *syscall emulation* and *full system simulation*. The former will simulate running a C program through a system call interface that is run directly on the host system. True to its name, the latter loads a virtual machine image and simulates the entire system.

We needed to simulate our selected benchmarks on the simulator which required a running JVM. As such, we took the full system simulation route. A known side effect of full system simulation is the large overhead required simulate *all* aspects of the virtual machine: logging into our simulated virtual machine took 1.5 hrs. We found that irrespective of simulation time for our benchmarks, the full system simulation made iterative development very challenging. We are aware that efforts to speed up simulation in general help to address this.

3 CHAPPIE Experimental Results

We include figures for all experiments for each of our benchmarks.




Figure 2. Top 10 Energy-Consuming Abstractions for batik




Figure 3. 1-CFA for batik

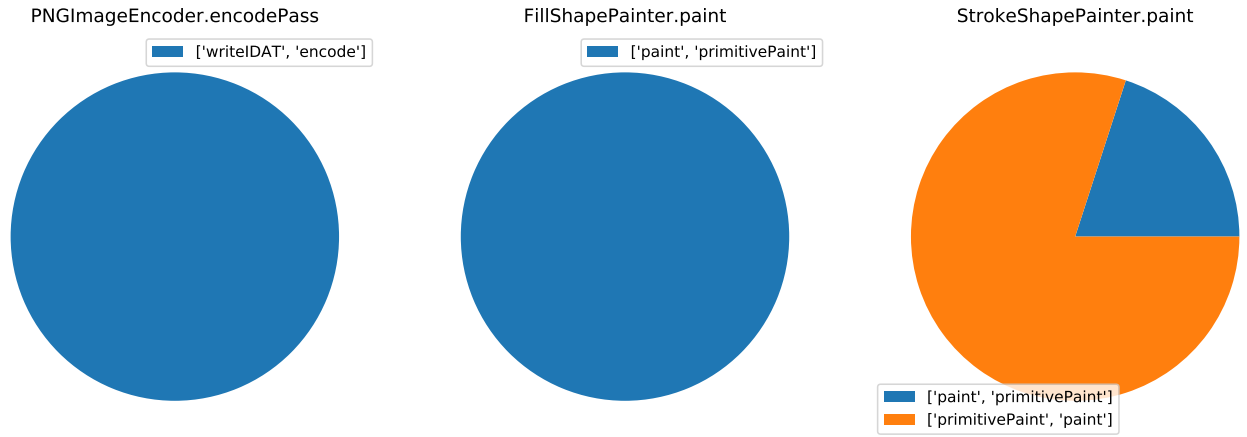


Figure 4. 2-CFA for batik




Figure 5. Top 10 Energy-Consuming Abstractions for biojava




Figure 6. 1-CFA for biojava




Figure 7. 2-CFA for biojava


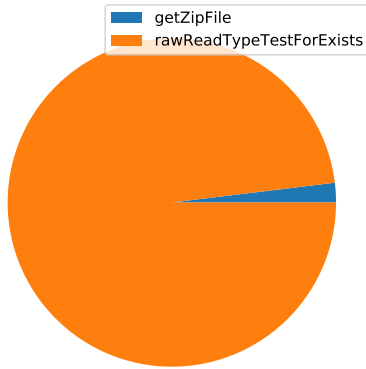
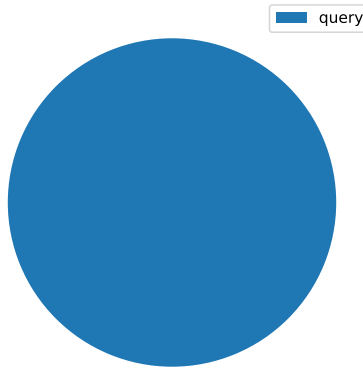


Figure 8. Top 10 Energy-Consuming Abstractions for eclipse

JavaModelManager.getZipFile



DiskIndex.addQueryResults



DiskIndex.readStreamChars

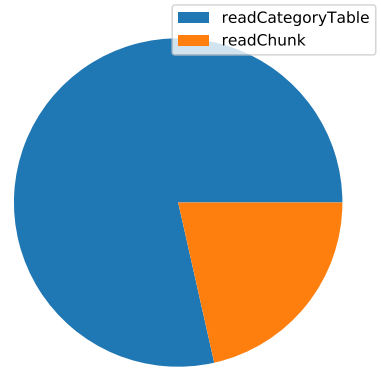




Figure 9. 1-CFA for eclipse

JavaModelManager.getZipFile



DiskIndex.addQueryResults



DiskIndex.readStreamChars




Figure 10. 2-CFA for eclipse




Figure 11. Top 10 Energy-Consuming Abstractions for fop




Figure 12. 1-CFA for fop




Figure 13. 2-CFA for fop




Figure 14. Top 10 Energy-Consuming Abstractions for graphchi




Figure 15. 1-CFA for graphchi




Figure 16. 2-CFA for graphchi




Figure 17. Top 10 Energy-Consuming Abstractions for h2




Figure 18. 1-CFA for h2




Figure 19. 2-CFA for h2




Figure 20. Top 10 Energy-Consuming Abstractions for jme




Figure 21. 1-CFA for jme



Figure 22. 2-CFA for jme




Figure 23. Top 10 Energy-Consuming Abstractions for jython




Figure 24. 1-CFA for jython




Figure 25. 2-CFA for jython




Figure 26. Top 10 Energy-Consuming Abstractions for lusearch




Figure 27. 1-CFA for lusearch




Figure 28. 2-CFA for lusearch

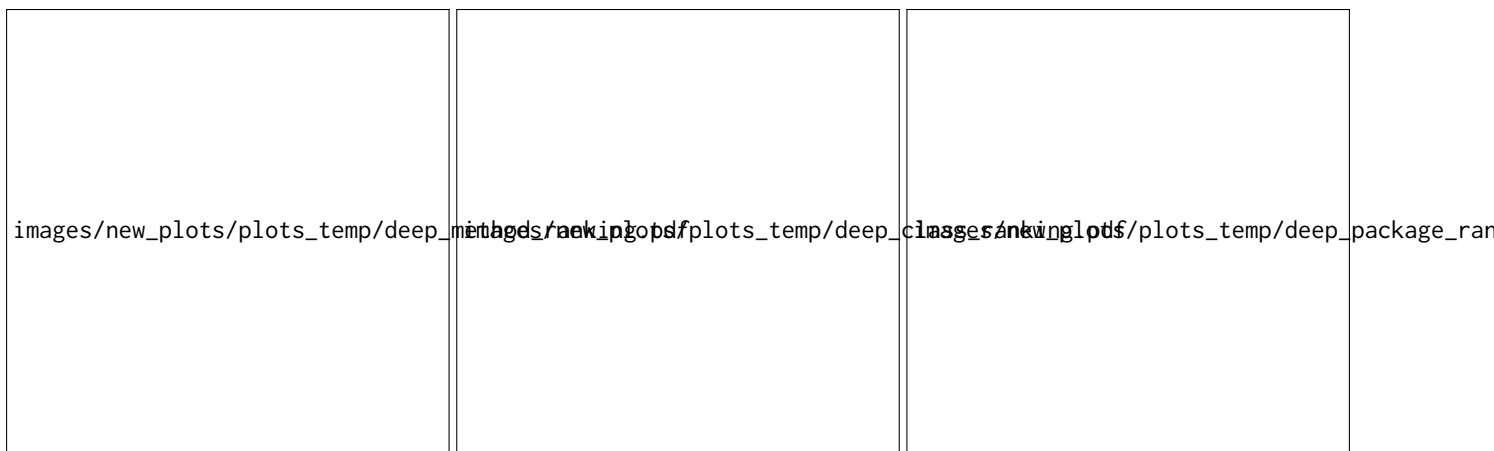


Figure 29. Top 10 Energy-Consuming Abstractions for plots_{temp}




Figure 30. 1-CFA for plots_{temp}




Figure 31. 2-CFA for plots_{temp}






Figure 32. Top 10 Energy-Consuming Abstractions for pmd

SourceCodeProcessor.usesTypeResolution



AbstractDelegateRule.apply



TeeOutputStream.write






Figure 33. 1-CFA for pmd

SourceCodeProcessor.usesTypeResolution



AbstractDelegateRule.apply



TeeOutputStream.write




Figure 34. 2-CFA for pmd




Figure 35. Top 10 Energy-Consuming Abstractions for sunflow




Figure 36. 1-CFA for sunflow




Figure 37. 2-CFA for sunflow




Figure 38. Top 10 Energy-Consuming Abstractions for xalan




Figure 39. 1-CFA for xalan




Figure 40. 2-CFA for xalan

References

- [1] BINKERT, N., BECKMANN, B., BLACK, G., REINHARDT, S. K., SAIDI, A., BASU, A., HESTNESS, J., HOWER, D. R., KRISHNA, T., SARDASHTI, S., SEN, R., SEWELL, K., SHOAIB, M., VAISH, N., HILL, M. D., AND WOOD, D. A. The gem5 simulator. *SIGARCH Comput. Archit. News* 39, 2 (Aug. 2011), 1–7.
- [2] LI, S., AHN, J. H., STRONG, R. D., BROCKMAN, J. B., TULSEN, D. M., AND JOUPPI, N. P. Mcpat: An integrated power, area, and timing modeling framework for multicore and manycore architectures. In *Proceedings of the 42Nd Annual IEEE/ACM International Symposium on Microarchitecture* (New York, NY, USA, 2009), MICRO 42, ACM, pp. 469–480.
- [3] SANDBERG, A., NIKOLERIS, N., CARLSON, T. E., HAGERSTEN, E., KAXIRAS, S., AND BLACK-SCHAFER, D. Full speed ahead: Detailed architectural simulation at near-native speed. In *2015 IEEE International Symposium on Workload Characterization* (Oct 2015), pp. 183–192.